

Machine Learning

As with the previous section, we'll look at some, but not all, popular machine learning packages in Julia. This part of the ecosystem is under heavy active development currently and could use more full-Julia options.

```
In [1]: using GaussianProcesses
import Random

Random.seed!(20140430)
# Training data
n=10; #number of training points
x = 2 $\pi$  * rand(n); #predictors
y = sin.(x) + 0.05*randn(n); #regressors
```

```

In [2]: #Select mean and covariance function
mZero = MeanZero() #Zero mean function
kern = SE(0.0,0.0) #Squared exponential kernel (note that hyperparameters are on the log scale)

logObsNoise = -1.0 # log standard deviation of observation noise (this is optional)
gp = GP(x,y,mZero,kern,logObsNoise) #Fit the GP

```

```

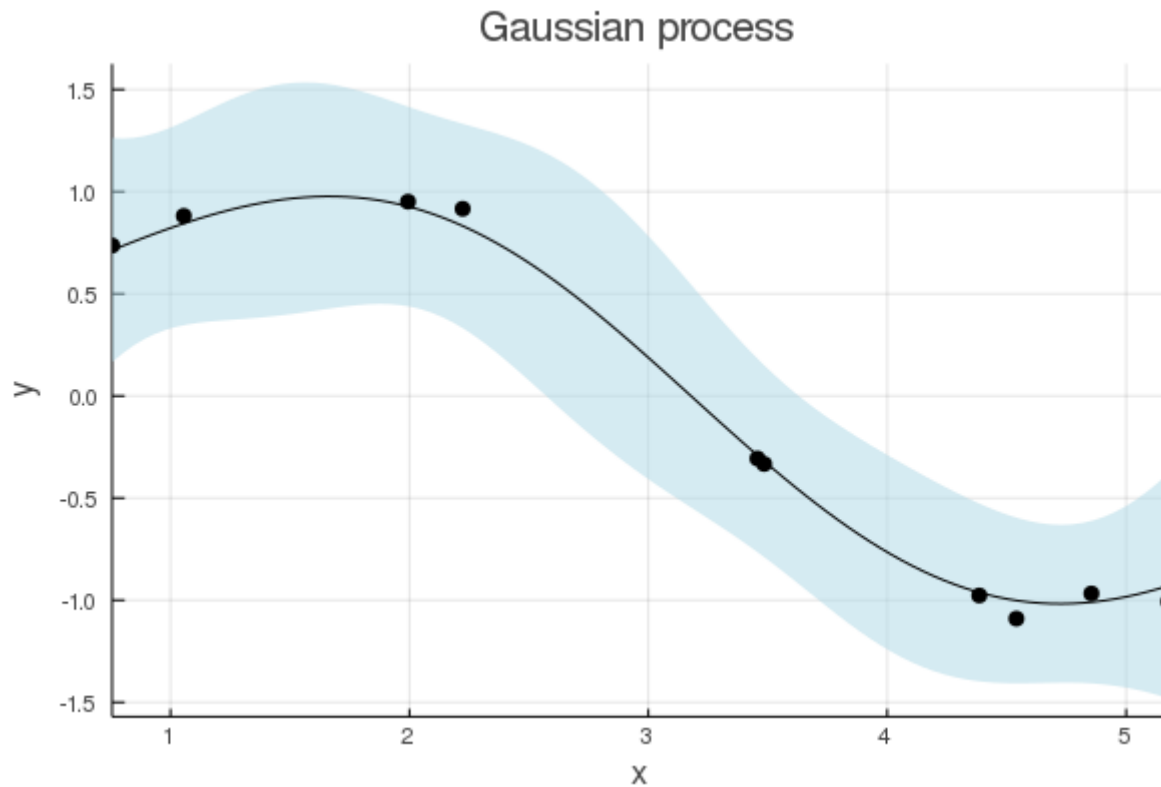
Out[2]: GP Exact object:
  Dim = 1
  Number of observations = 10
  Mean function:
    Type: MeanZero, Params: Float64[]
  Kernel:
    Type: SEIso{Float64}, Params: [0.0, 0.0]
  Input observations =
[4.854610892030431 5.176527683588912 ... 1.9941157865846477 3.4567590405289352]
  Output observations = [-0.9672931901680022, -1.0070469610508248, -1.09039628
29715461, 0.8811208605233783, -0.3332131916909575, -0.9769651482106089, 0.9159
341732111771, 0.7362180710829012, 0.9508490881368652, -0.3064317818816142]
  Variance of observation noise = 0.1353352832366127
  Marginal Log-Likelihood = -6.335

```

In [3]: `import Plots`

```
Plots.plot(gp; xlabel="x", ylabel="y", title="Gaussian process", legend=false, f
mt=:png)
```

Out[3]:



```
In [4]: using Optim
        rosenbrock(x) = (1.0 - x[1])^2 + 100.0 * (x[2] - x[1]^2)^2
        result = optimize(rosenbrock, zeros(2), BFGS())
```

```
Out[4]: * Status: success

* Candidate solution
  Minimizer: [1.00e+00, 1.00e+00]
  Minimum:   5.471433e-17

* Found with
  Algorithm:      BFGS
  Initial Point: [0.00e+00, 0.00e+00]

* Convergence measures
  |x - x'|          = 3.47e-07  $\nless$  0.0e+00
  |x - x'|/|x'|     = 3.47e-07  $\nless$  0.0e+00
  |f(x) - f(x')|    = 6.59e-14  $\nless$  0.0e+00
  |f(x) - f(x')|/|f(x')| = 1.20e+03  $\nless$  0.0e+00
  |g(x)|            = 2.33e-09  $\leq$  1.0e-08

* Work counters
  Seconds run:      1 (vs limit Inf)
  Iterations:       16
  f(x) calls:       53
   $\nabla$ f(x) calls:    53
```

```
In [5]: optimize!(gp; method=ConjugateGradient())
```

```
Out[5]: * Status: success

* Candidate solution
  Minimizer: [-2.99e+00, 4.64e-01, -2.29e-01]
  Minimum:   -3.275745e+00


* Found with
  Algorithm:   Conjugate Gradient
  Initial Point: [-1.00e+00, 0.00e+00, 0.00e+00]

* Convergence measures
  |x - x'|           = 2.12e-09  $\leq$  0.0e+00
  |x - x'|/|x'|      = 7.08e-10  $\leq$  0.0e+00
  |f(x) - f(x')|     = 1.42e-14  $\leq$  0.0e+00
  |f(x) - f(x')|/|f(x')| = 4.34e-15  $\leq$  0.0e+00
  |g(x)|             = 3.80e-09  $\leq$  1.0e-08

* Work counters
  Seconds run:   0 (vs limit Inf)
  Iterations:   26
  f(x) calls:   64
   $\nabla$ f(x) calls: 40
```

```
In [6]: Plots.plot(gp; legend=false, fmt=:png)
```

Out[6]:



DecisionTree.jl

Provides Decision Trees and Random Forests, with a scikit-learn based API

In [7]: `using DecisionTree`


```
In [8]: using DataFrames, RDatasets
iris = dataset("datasets", "iris");
features = convert(Array, iris[:, 1:4])
labels = string.(iris[:, :Species])
features, labels
```

```
Out[8]: ([5.1 3.5 1.4 0.2; 4.9 3.0 1.4 0.2; ... ; 6.2 3.4 5.4 2.3; 5.9 3.0 5.1 1.8], ["s
etosa", "setosa", "setosa", "setosa", "setosa", "setosa", "setosa", "setosa", "setosa",
"setosa", "setosa" ... "virginica", "virginica", "virginica", "virginica", "vi
rginica", "virginica", "virginica", "virginica", "virginica", "virginica"])
```

```
In [9]: model = DecisionTreeClassifier(max_depth=2)
```

```
Out[9]: DecisionTreeClassifier
max_depth:          2
min_samples_leaf:   1
min_samples_split:  2
min_purity_increase: 0.0
pruning_purity_threshold: 1.0
n_subfeatures:      0
classes:            nothing
root:               nothing
```

```
In [10]: DecisionTree.fit!(model, features, labels)
print_tree(model)
```

```
Feature 3, Threshold 2.45
L-> setosa : 50/50
R-> Feature 4, Threshold 1.75
    L-> versicolor : 49/54
    R-> virginica : 45/46
```

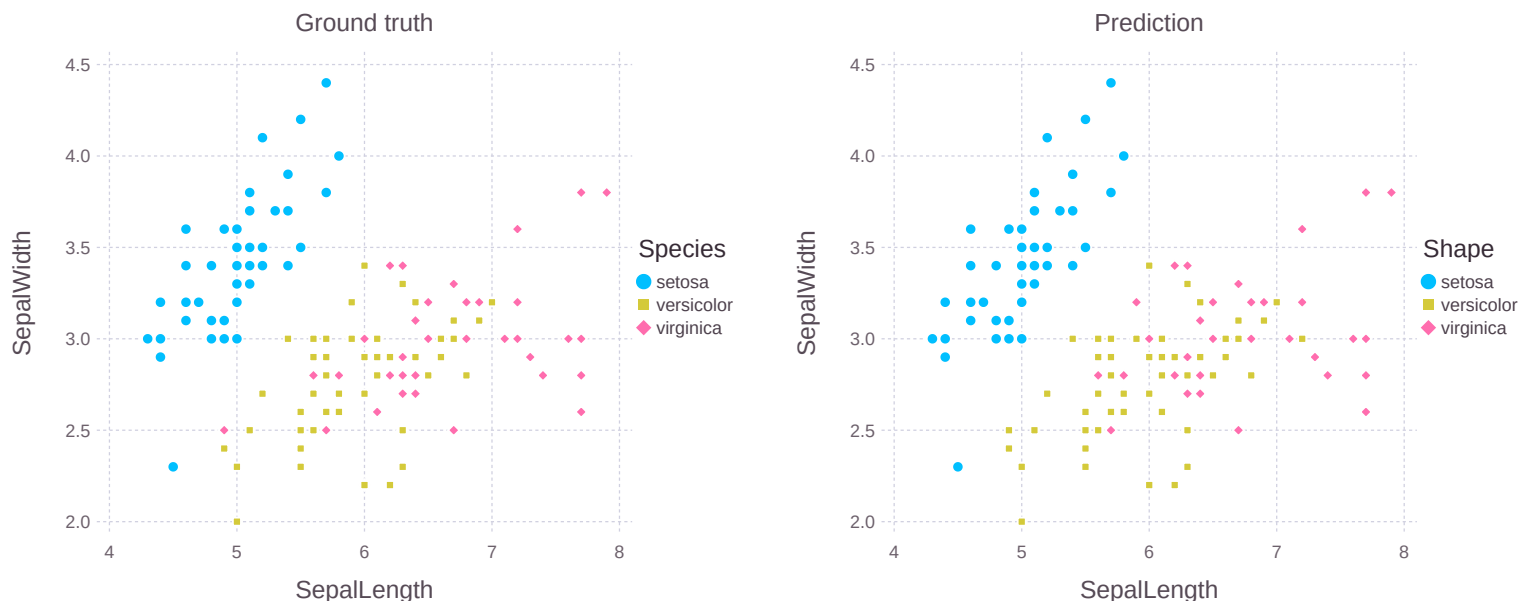
```
In [11]: y = DecisionTree.predict(model, features);  
         sum(y .!= iris[:, :Species])
```

```
Out[11]: 6
```

```
In [12]: using Gadfly
p1 = plot(iris, x="SepalLength", y="SepalWidth", color="Species", shape="Species",
          Guide.title("Ground truth"));
p2 = plot(iris, x="SepalLength", y="SepalWidth", color=y, shape=y, Geom.point,
          Guide.title("Prediction"));
set_default_plot_size(700pt, 300pt)
hstack(p1, p2)
```

```
[ Info: Loading DataFrames support into Gadfly.jl
  @ Gadfly /home/d9w/.julia/packages/Gadfly/09PWZ/src/mapping.jl:228
[ Warning: `getindex(df::DataFrame, col_ind::ColumnIndex)` is deprecated, use
`df[!, col_ind]` instead.
  caller = evalmapping at dataframes.jl:96 [inlined]
  @ Core /home/d9w/.julia/packages/Gadfly/09PWZ/src/dataframes.jl:96
```

Out[12]:



ScikitLearn.jl

Uses PyCall to adapt the ScikitLearn library to Julia and adds interaction with existing Julia libraries, like GaussianProcesses.jl and DecisionTrees.jl

```
In [13]: using ScikitLearn
@sk_import linear_model: LogisticRegression
model = LogisticRegression(fit_intercept=true)
```

```
Out[13]: PyObject LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercep
t=True,
                                     intercept_scaling=1, l1_ratio=None, max_iter=100,
                                     multi_class='auto', n_jobs=None, penalty='l2',
                                     random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                                     warm_start=False)
```

```
In [14]: X = convert(Array, iris[!, [:SepalLength, :SepalWidth, :PetalLength, :PetalWidth
]])
y = convert(Array, iris[!, :Species])
fit!(model, X, y)

accuracy = sum(ScikitLearn.predict(model, X) .== y) / length(y)
println("accuracy: $accuracy")
```

```
accuracy: 0.9733333333333334
```

Flux.jl

A general machine learning library with a focus on deep learning, entirely in Julia

```
In [15]: using Flux, Statistics
using Flux: onehotbatch, onecold, crossentropy, throttle
using Base.Iterators: repeated
using Flux.Data.FashionMNIST
```


```
└ Info: CUDAdrv.jl failed to initialize, GPU functionality unavailable (set JULIA_CUDA_SILENT or JULIA_CUDA_VERBOSE to silence or expand this message)
└ @ CUDAdrv /home/d9w/.julia/packages/CUDAdrv/mCr00/src/CUDAdrv.jl:69
```

```
In [16]: train_imgs = FashionMNIST.images(:train)
train_labels = FashionMNIST.labels(:train)
text_labels = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat", "Sandal",
               "Shirt", "Sneaker", "Bag", "Ankle boot"];
```

```
In [17]: println(text_labels[train_labels[1]+1])  
train_imgs[1]
```

Ankle boot

Out[17]:




```
In [18]: batch_size = 128

X = hcat(float.(reshape.(train_imgs, :)))...
Y = onehotbatch(train_labels, 0:9)
batches=[(X[:,i],Y[:,i]) for i in Iterators.partition(1:length(train_imgs),batch_size)];
```

```
In [19]: SimpleNet = Chain(  
          Dense(28^2,120),  
          Dense(120,10),softmax)
```

```
Out[19]: Chain(Dense(784, 120), Dense(120, 10), softmax)
```

```
In [20]: loss(x,y) = crossentropy(SimpleNet(x),y)  
         opt = Descent()  
  
         nn_accuracy(x, y) = mean(onecold(cpu(SimpleNet(x))) .== onecold(cpu(y)))  
         evalcb = throttle(()->println(nn_accuracy(X,Y)), 1);
```

```
In [21]: Flux.train!(loss, params(SimpleNet), batches, opt, cb = evalcb)
```

```
0.1498  
0.7043  
0.7396333333333334  
0.7269833333333333  
0.7887666666666666  
0.79895  
0.80735  
0.8014166666666667  
0.8012666666666667  
0.8209166666666666  
0.8253  
0.8227833333333333  
0.8195  
0.8290833333333333  
0.8320166666666666  
0.82685  
0.8155  
0.79055
```

```
In [22]: test_imgs = FashionMNIST.images(:test)
         predictions = onecold(SimpleNet(hcat(float.(reshape.(test_imgs, :))....)))
         test_labels = FashionMNIST.labels(:test) .+ 1;
```

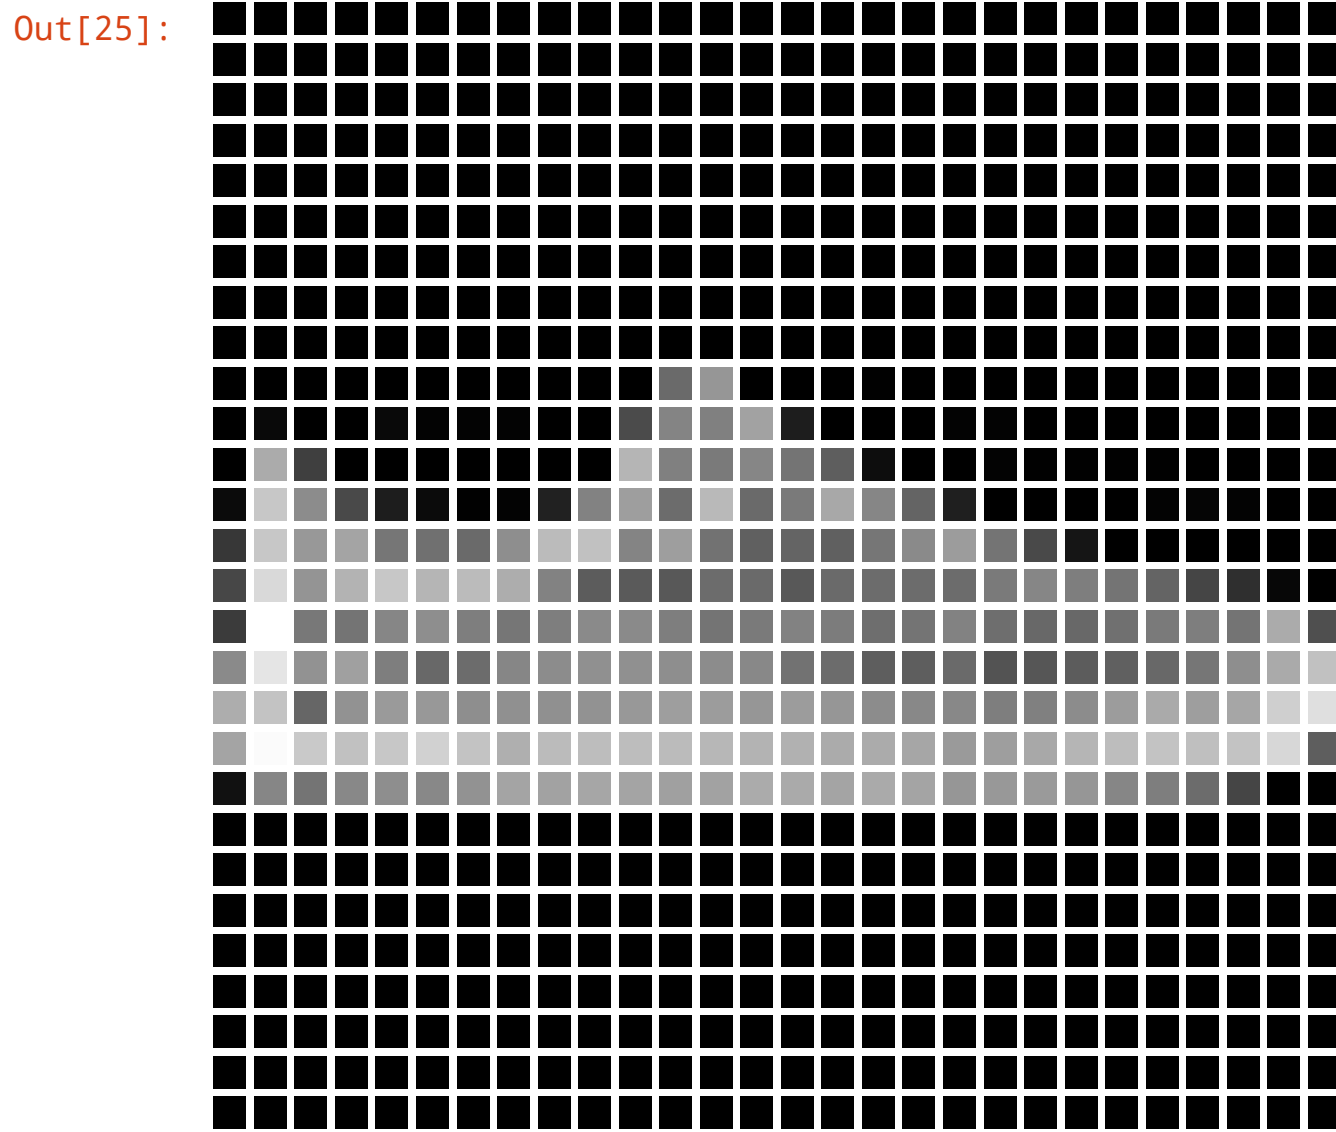
```
In [23]: sum(predictions .== test_labels) / length(test_labels)
```

```
Out[23]: 0.8236
```

```
In [24]: wrong = predictions .!= test_labels
println("Prediction\tReal Label")
for i in findall(wrong)[1:5]
    println(text_labels[predictions[i]], "\t\t",
            text_labels[test_labels[i]])
end
```

Prediction	Real Label
Sandal	Sneaker
T-shirt/top	Pullover
Sneaker	Sandal
Sneaker	Ankle boot
Pullover	Coat

```
In [25]: test_imgs[findfirst(wrong)]
```



```
In [ ]:
```